

# M. Farrukh Mehmood

Machine Learning Engineer | Industrial AI

**Portfolio:** [sfarrukhm.github.io](https://sfarrukhm.github.io)

**LinkedIn:** [in/sfarrukhm](https://www.linkedin.com/in/sfarrukhm)

**Contact:** 03017569039

**Email:** [smfarrukhm@gmail.com](mailto:smfarrukhm@gmail.com)

**Address:** Bahawalpur, Pakistan

---

## Technical Skills & Tools

---

- **Machine Learning:** Supervised and unsupervised learning, time series forecasting, clustering, PCA, feature engineering, explainable AI
- **Computer Vision:** OpenCV, CNNs, object detection, image segmentation, image classification, feature extraction, augmentation, transfer learning
- **Agentic AI:** LangChain, LangGraph, tool calling, prompt engineering, RAG, OpenAI API, Gemini CLI
- **Deployment & MLOps:** Docker, FastAPI, ONNX, GitHub Actions, REST APIs, MLflow, CI/CD, model serialization
- **Languages & Tools:** Python, C/C++, SQL, Bash, Git
- **ML Frameworks:** PyTorch, TensorFlow, Scikit-learn, XGBoost, Hugging Face Transformers
- **Databases:** MongoDB, PostgreSQL, SQLite, vector databases (ChromaDB, Pinecone)
- **Web App Development:** Streamlit, Gradio

## RELEVANT PROJECTS

---

### Model Compression for Edge Deployment

#### *Intel Image Classification with Knowledge Distillation*

- Implemented a teacher-student distillation pipeline to compress a ResNet18-based spatial attention model into a lightweight CNN.
- Trained the student model using both ground-truth labels and soft targets from the teacher to preserve accuracy.
- Achieved 99.6% reduction in parameters and a 4× improvement in inference speed with minimal performance loss.
- Analyzed confusion matrices and output snapshots to evaluate error consistency and classification behavior.

**Tools/Techniques:** PyTorch, ResNet18, Knowledge Distillation, Attention, Model Optimization

### Sensor Data Processing & Time-Series Modeling

#### *CMI - Detect Behavior with Sensor Data*

Ongoing [Kaggle]

- Engineered statistical and domain-informed features from high-frequency motion and temperature sensor data.
- Built and optimized a LightGBM classification pipeline achieving a **CV score of 0.74** (mean of binary and macro F1).
- Applied signal segmentation, feature extraction, and gradient boosting to transform time-series into actionable tabular insights.

**Tools/Techniques:** LightGBM, Pandas, scikit-learn, Time-series feature engineering

## Agentic AI for Cybersecurity

### *Autonomous Threat Intelligence Reporter*

- Developed an agent-based app that inspects API endpoints for potential threats using platforms like VirusTotal.
- Used LangGraph to orchestrate tool use and Gemini API to generate human-readable threat reports.
- Built a Streamlit interface for input submission and report visualization.

**Tools/Techniques:** LangGraph, Gemini API, Streamlit, Agentic AI

## Microservice-Based Digital Platform

### *RUL Inference Engine*

- Built a predictive system for Remaining Useful Life (RUL) estimation using NASA's C-MAPSS turbofan engine dataset.
- Trained both XGBoost and LSTM models to capture degradation trends.
- Deployed preprocessing and inference modules as FastAPI microservices, exposed via REST APIs.
- Containerized the system with Docker for scalable, environment-agnostic deployment..

**Tools/Techniques:** FastAPI, Docker, XGBoost, LSTM

## Deployable ML Application with User Interface

### *Amazon Listing Optimization Tool*

- Developed an LLM-powered interface enabling users to generate optimized Amazon product listings from minimal product input.
- Designed a no-code user experience where form inputs triggered a keyword extraction and content generation pipeline.

**Tools/Techniques:** Streamlit, Hugging Face Transformers, RAG, Polars

## WORK EXPERIENCE

---

Research Assistant — SMME, NUST

Dec 2023 – Dec 2024

### *Generative Modeling of Unsteady Flow Fields*

- Developed a generative sequence model combining Variational Autoencoder (VAE) and LSTM for temporal flow field prediction.
- Trained on CFD simulation data to predict future flow snapshots from a few initial inputs, achieving 95% accuracy.
- Contributed to research on applying ML techniques to CFD-generated data for modeling unsteady flow dynamics.

**Tools/Techniques:** VAE, LSTM, PyTorch, CFD, Time Series Modeling

## WORKSHOP DELIVERED

---

### **Predicting Particle Deposition in Lung Airways Using the Discrete Phase Model (DPM)**

SMME, NUST | Oct. 2024

Delivered a hands-on workshop covering theoretical and practical aspects of particle deposition modeling using DPM. Included simulation setup, live CFD demonstrations, and interpretation of health-related aerosol transport results.

## Relevant Courses

---

- Machine Learning — SEECS, NUST
- Deep Learning — SMME, NUST
- Machine Learning in Production — Coursera
- MLOps | Machine Learning Operations Specialization — Duke University
- IT Automation with Python — Google Professional Certificate
- NLP with Classification and Vector Spaces — Coursera

## EDUCATION

---

### MS in Mechanical Engineering

2022-2024

NUST, Islamabad, Pakistan

CGPA: 3.86/4.00

Thesis Title: *Numerical Investigation of Aerosol Deposition in Lung Airways During Inhalation and Exhalation*

### BSc. in Mechanical Engineering

2014-2018

UET, Taxila, Pakistan

## HONORS, VOLUNTEERING & LANGUAGES

---

- **Awards:**
  - Best Project — Machine Learning, SEECS NUST (2023)
  - Water Rocket Champion — UET Taxila (2015)
  - PEEF Scholarship — Awarded throughout undergraduate studies (2014 – 2018)
- **Volunteering:**
  - Organizing Member – ASME UET Chapter
  - Finance Head – Umeed-e-Subh Welfare Society
- **Languages:** English, Urdu